


Experiment 1 : Data Cleaning & EDA

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import warnings
```


```
1 data = pd.read_csv("/content/data.csv")
```

```
1 data.head(5)
```



	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	NaN	Yes


```
1 data.shape
```



(10, 4)


```
1 pd.set_option('display.max_columns', None)
2 pd.set_option('display.max_rows', None)
```

```
1 data.head(2)
```




	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes

```
1 data.tail(2)
```



	Country	Age	Salary	Purchased
8	Germany	50.0	83000.0	No
9	France	37.0	67000.0	Yes


```
1 data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Country     10 non-null     object
 1   Age         9 non-null      float64
 2   Salary      9 non-null      float64
 3   Purchased   10 non-null     object
```

dtypes: float64(2), object(2)  
memory usage: 448.0+ bytes

```
1 data.isnull().sum()
```




	0
Country	0
Age	1
Salary	1
Purchased	0

dtype: int64

```
1 plt.figure(figsize=(25,25))
2 sns.heatmap(data.isnull())
3 plt.show()
```



0  
1  
2  
3  
4  
5  
6



1.0  
0.8  
0.6  
0.4

7

0.2

```
1 missing_value_percent = data.isnull().sum() / data.shape[0] * 100
2 print(missing_value_percent)
```

```
↵ Country      0.0
   Age         10.0
   Salary      10.0
   Purchased    0.0
dtype: float64
```

```
1 missing_value_column = missing_value_percent[missing_value_percent > 17].keys()
2 print(missing_value_column)
```

```
↵ Index([], dtype='object')
```

```
1 data1 = data.drop(columns = missing_value_column)
```

```
1 data1.shape
```


```
↵ (10, 4)
```

```
1 data2 = data1.dropna()
```

```
1 data2.shape
```

```
↵ (8, 4)
```

```
1 plt.figure(figsize=(25,25))
2 sns.heatmap(data1.isnull())
3 plt.show()
```





```
1 data2.isnull().sum().sum()
```

→ 0

Experiment 2: Outliers estimation & Plot it through violen/whisker

```
1 # Importing
2 import sklearn
3 from sklearn.datasets import load_diabetes
4 import pandas as pd
5 import matplotlib.pyplot as plt
6
7 # Load the dataset
8 diabetes = load_diabetes()
9
10 # Create the dataframe
11 column_name = diabetes.feature_names
12 df_diabetes = pd.DataFrame(diabetes.data)
13 df_diabetes.columns = column_name
14 print(df_diabetes.head())
15
```

→

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	

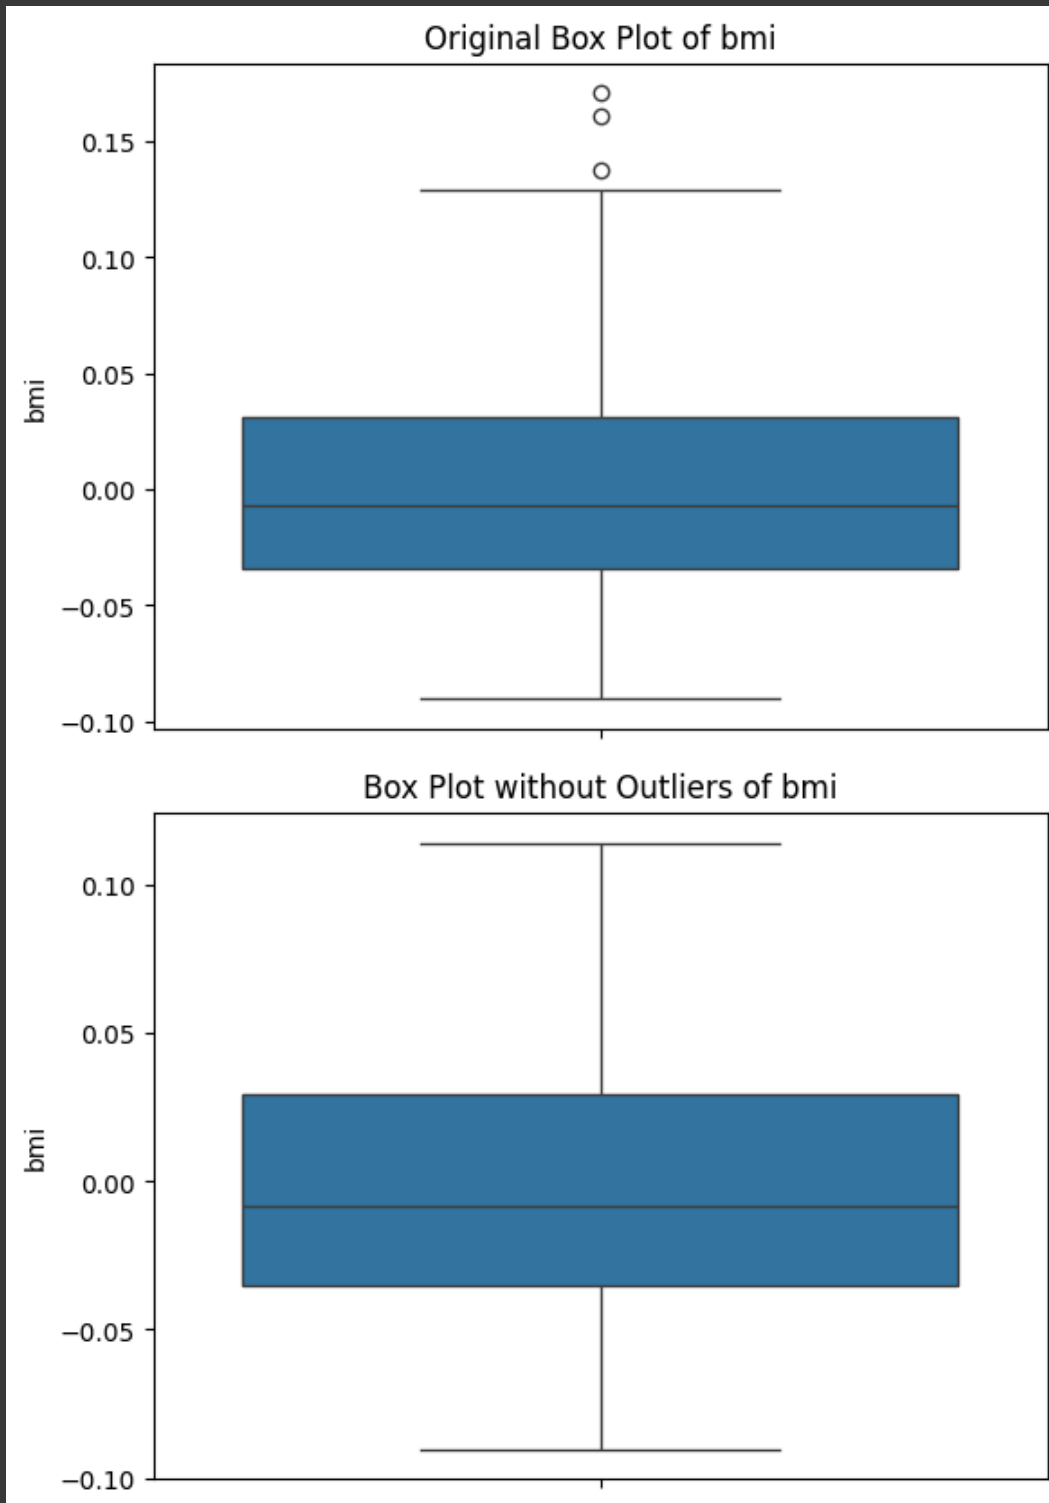
	s4	s5	s6
0	-0.002592	0.019907	-0.017646
1	-0.039493	-0.068332	-0.092204
2	-0.002592	0.002861	-0.025930
3	0.034309	0.022688	-0.009362
4	-0.002592	-0.031988	-0.046641

```
1 import seaborn as sns
2 import matplotlib.pyplot as plt
3
4
5 def removal_box_plot(df, column, threshold):
6     sns.boxplot(df[column])
7     plt.title(f'Original Box Plot of {column}')
8     plt.show()
9
10     removed_outliers = df[df[column] <= threshold]
11
12     sns.boxplot(removed_outliers[column])
13     plt.title(f'Box Plot without Outliers of {column}')
14     plt.show()
15     return removed_outliers
```

```

16
17
18 threshold_value = 0.12
19
20 no_outliers = removal_box_plot(df_diabetics, 'bmi', threshold_value)
21

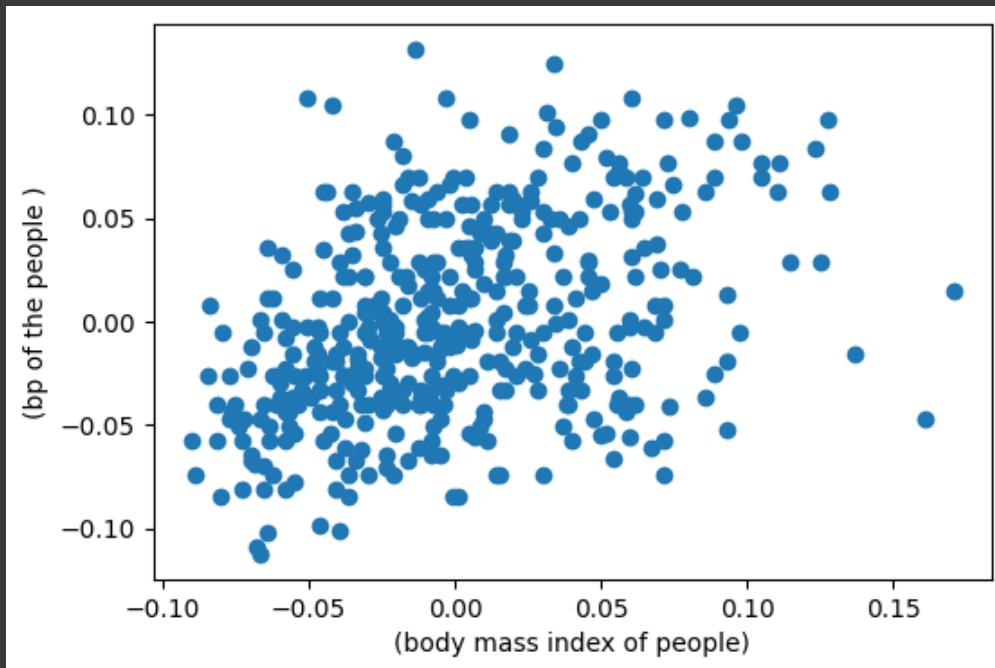
```




```

1 fig, ax = plt.subplots(figsize=(6, 4))
2 ax.scatter(df_diabetics['bmi'], df_diabetics['bp'])
3 ax.set_xlabel('(body mass index of people)')
4 ax.set_ylabel('(bp of the people )')
5 plt.show()
6

```



```
1 import numpy as np
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4
5 outlier_indices = np.where((df_diabetics['bmi'] > 0.12) &
6 (df_diabetics['bp'] < 0.8))
7
8 no_outliers = df_diabetics.drop(outlier_indices[0])
9
10 # Scatter plot without outliers
11 fig, ax_no_outliers = plt.subplots(figsize=(6, 4))
12 ax_no_outliers.scatter(no_outliers['bmi'], no_outliers['bp'])
13 ax_no_outliers.set_xlabel('(body mass index of people)')
14 ax_no_outliers.set_ylabel('(bp of the people )')
15 plt.show()
16
```



```
1 import numpy as np
2
3 threshold_z = 2
```



```
4
5 outlier_indices = np.where(z > threshold_z)[0]
6 no_outliers = df_diabetics.drop(outlier_indices)
7 print("Original DataFrame Shape:", df_diabetics.shape)
8 print("DataFrame Shape after Removing Outliers:", no_outliers.shape)
9
```

```
↔ Original DataFrame Shape: (442, 10)
   DataFrame Shape after Removing Outliers: (426, 10)
```

```
1 # IQR
2 Q1 = np.percentile(df_diabetics['bmi'], 25, method='midpoint')
3 Q3 = np.percentile(df_diabetics['bmi'], 75, method='midpoint')
4 IQR = Q3 - Q1
5 print(IQR)
6
```

```
↔ 0.06520763046978838
```

```
1 # Above Upper bound
2 upper = Q3+1.5*IQR
3 upper_array = np.array(df_diabetics['bmi'] >= upper)
4 print("Upper Bound:", upper)
5 print(upper_array.sum())
6
7 # Below Lower bound
8 lower = Q1-1.5*IQR
9 lower_array = np.array(df_diabetics['bmi'] <= lower)
10 print("Lower Bound:", lower)
11 print(lower_array.sum())
12
```

```
↔ Upper Bound: 0.12879000811776306
   3
   Lower Bound: -0.13204051376139045
   0
```

```
1 # Importing
2 import sklearn
3 from sklearn.datasets import load_diabetes
4 import pandas as pd
5
6 # Load the dataset
7 diabetes = load_diabetes()
8
9 # Create the dataframe
10 column_name = diabetes.feature_names
11 df_diabetes = pd.DataFrame(diabetes.data)
12 df_diabetes.columns = column_name
13 df_diabetes.head()
14 print("Old Shape: ", df_diabetes.shape)
15
16 ''' Detection '''
17 # IQR
18 # Calculate the upper and lower limits
19 Q1 = df_diabetes['bmi'].quantile(0.25)
20 Q3 = df_diabetes['bmi'].quantile(0.75)
21 IQR = Q3 - Q1
22 lower = Q1 - 1.5*IQR
23 upper = Q3 + 1.5*IQR
24
25 # Create arrays of Boolean values indicating the outlier rows
26 upper_array = np.where(df_diabetes['bmi'] >= upper)[0]
```

```
27 lower_array = np.where(df_diabetes['bmi'] <= lower)[0]
28
29 # Removing the outliers
30 df_diabetes.drop(index=upper_array, inplace=True)
31 df_diabetes.drop(index=lower_array, inplace=True)
32
33 # Print the new shape of the DataFrame
34 print("New Shape: ", df_diabetes.shape)
35
```

```
➦ Old Shape: (442, 10)
New Shape: (439, 10)
```

Experiment 3: Imputation of Missing Values KNN, Interpolation,Min-Max/Z-score.

```
1 # import necessary libraries
2 import numpy as np
3 import pandas as pd
4
5 # import the KNNImputer class
6 from sklearn.impute import KNNImputer
7
8
9 # create dataset for marks of a student
10 dict = {'Maths': [80, 90, np.nan, 95],
11         'Chemistry': [60, 65, 56, np.nan],
12         'Physics': [np.nan, 57, 80, 78],
13         'Biology': [78, 83, 67, np.nan]}
14
15 # creating a data frame from the list
16 Before_imputation = pd.DataFrame(dict)
17 # print dataset before imputation
18 print("Data Before performing imputation\n", Before_imputation)
19
20 # create an object for KNNImputer
21 imputer = KNNImputer(n_neighbors=2)
22 After_imputation = imputer.fit_transform(Before_imputation)
23 # print dataset after performing the operation
24 print("\n\nAfter performing imputation\n", After_imputation)
25
```

```
➦ Data Before performing imputation
   Maths  Chemistry  Physics  Biology
0   80.0         60.0      NaN    78.0
1   90.0         65.0     57.0    83.0
2    NaN         56.0     80.0    67.0
3   95.0          NaN     78.0     NaN
```

```
After performing imputation
[[80.  60.  68.5 78. ]
 [90.  65.  57.  83. ]
 [87.5 56.  80.  67. ]
 [95.  58.  78.  72.5]]
```


```
1 import pandas as pd
2 import numpy as np
3 a=pd.Series([0, 1, np.nan, 3,4,5,7])
4 a.interpolate()
```





```
<ipython-input-3-14e7dc7c7159>:1: FutureWarning: The default value of observed=False is deprecated and will change to observed=True in a future version of pandas. Specify observed=False to silence this warning and r
pivot = np.round(pd.pivot_table(data, values='price',
<ipython-input-3-14e7dc7c7159>:1: FutureWarning: The provided callable <function mean at 0x7f58e7b6c280> is currently using DataFrameGroupBy.mean. In a future version of pandas, the provided callable will be used di
pivot = np.round(pd.pivot_table(data, values='price',
```

	fuel-type	diesel	gas							
	fuel-system	idi	1bb1	2bb1	4bb1	mfi	mpfi	spdi	spfi	
num-of-doors	body-style									
four	hatchback	7788.00	0.00	7813.71	0.0	0.0	10618.00	0.00	0.0	
	sedan	16328.92	8811.67	7711.19	0.0	0.0	18425.68	9279.00	0.0	
	wagon	19727.67	7295.00	8028.89	0.0	0.0	14213.42	0.00	0.0	
two	convertible	0.00	0.00	0.00	0.0	0.0	21890.50	0.00	0.0	
	hardtop	28176.00	0.00	8249.00	0.0	0.0	23540.50	0.00	0.0	
	hatchback	0.00	7054.43	6701.67	12145.0	12964.0	14581.50	11479.43	11048.0	
	sedan	7437.00	0.00	7570.00	0.0	0.0	21034.00	0.00	0.0	




```
1 np.round(pd.pivot_table(data, values='price',
2                               index=['body-style'],
3                               columns=['num-of-doors'],
4                               aggfunc=[np.mean, np.median],
5                               fill_value=0),2)
```



```
<ipython-input-4-18854bd411ba>:1: FutureWarning: The default value of observed=False is deprecated and will change to observed=True in a future version of pandas. Specify observed=False to silence this warning and r
np.round(pd.pivot_table(data, values='price',
<ipython-input-4-18854bd411ba>:1: FutureWarning: The provided callable <function mean at 0x7f58e7b6c280> is currently using DataFrameGroupBy.mean. In a future version of pandas, the provided callable will be used di
np.round(pd.pivot_table(data, values='price',
<ipython-input-4-18854bd411ba>:1: FutureWarning: The default value of observed=False is deprecated and will change to observed=True in a future version of pandas. Specify observed=False to silence this warning and r
np.round(pd.pivot_table(data, values='price',
<ipython-input-4-18854bd411ba>:1: FutureWarning: The provided callable <function median at 0x7f58e755b910> is currently using DataFrameGroupBy.median. In a future version of pandas, the provided callable will be use
np.round(pd.pivot_table(data, values='price',
```

		mean		median	
num-of-doors		four	two	four	two
body-style					
convertible		0.00	21890.50	0.0	17084.5
	hardtop	0.00	22208.50	0.0	19687.5
hatchback		8372.00	10230.79	8073.0	8970.0
sedan		14614.13	14283.00	12555.0	8678.0
wagon		12371.96	0.00	11694.0	0.0



```
1 import pandas as pd
2 d1 = {"Name": ["Tom", "Jerry", "Spike"], "ID": [1, 2, 3],
3       "Role": ["Cat", "Mouse", "Dog"]}
4 df = pd.DataFrame(d1)
5 print(df)
6
7 df_melted = pd.melt(df, id_vars=["ID"], value_vars=["Name", "Role"])
8 print(df_melted)
```



	Name	ID	Role
0	Tom	1	Cat

```

1 Jerry 2 Mouse
2 Spike 3 Dog
   ID variable value
0 1 Name Tom
1 2 Name Jerry
2 3 Name Spike
3 1 Role Cat
4 2 Role Mouse
5 3 Role Dog

```

```

1 #multiple columns as id_vars
2 df_melted = pd.melt(df, id_vars=["ID", "Name"], value_vars=["Role"])
3 print(df_melted)

```

```

⇌ ID Name variable value
0 1 Tom Role Cat
1 2 Jerry Role Mouse
2 3 Spike Role Dog

```

```

1 #skipping columns in melt function
2 df_melted = pd.melt(df, id_vars=["Name"], value_vars=["Role"])
3 print(df_melted)

```

```

⇌ Name variable value
0 Tom Role Cat
1 Jerry Role Mouse
2 Spike Role Dog

```

Experiment 5: Image normalization (rescale -1,1).

```

1 import numpy as np
2 from PIL import Image
3 import matplotlib.pyplot as plt
4
5 # Load an image
6 image_path = r"/content/mascot-logo-design_P1_900x420.jpg"
7
8 try:
9     image = Image.open(image_path)
10    image = image.convert('RGB') # Ensure image is in RGB format
11 except FileNotFoundError:
12    print(f"File not found: {image_path}")
13    raise
14 except Exception as e:
15    print(f"Error occurred: {e}")
16    raise
17
18 # Convert image to numpy array
19 image_array = np.array(image).astype('float32')
20
21 # Normalize the image to the range [-1, 1]
22 normalized_image_array = (image_array / 127.5) - 1
23
24 # Verify normalization
25 print(f"Min value in normalized image: {np.min(normalized_image_array)}")
26 print(f"Max value in normalized image: {np.max(normalized_image_array)}")
27
28 # Convert normalized image back to the range [0, 255] for visualization
29 denormalized_image_array = (normalized_image_array + 1) * 127.5
30 denormalized_image_array = denormalized_image_array.astype('uint8')
31
32 # Display the original and normalized images

```

```

33 fig, axes = plt.subplots(1, 2, figsize=(12, 6))
34
35 axes[0].imshow(image)
36 axes[0].set_title("Original Image")
37 axes[0].axis('off')
38
39 axes[1].imshow(denormalized_image_array)
40 axes[1].set_title("Normalized Image (rescaled to [0, 255] for display)")
41 axes[1].axis('off')
42
43 plt.show()
44

```



```

Min value in normalized image: -1.0
Max value in normalized image: 1.0

```

Original Image



Normalized Image (rescaled to [0, 255] for display)



Experiment 6: Linear regression and Logistic regression.

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 from sklearn.model_selection import train_test_split
5 from sklearn.linear_model import LinearRegression
6 from sklearn.metrics import mean_squared_error, r2_score
7
8 # Create a synthetic dataset for linear regression
9 np.random.seed(0)
10 X = 2 * np.random.rand(100, 1)
11 y = 4 + 3 * X + np.random.randn(100, 1)
12
13 # Split the dataset into training and testing sets
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
15                                                    random_state=0)
16
17 # Create and train the linear regression model
18 lin_reg = LinearRegression()
19 lin_reg.fit(X_train, y_train)
20
21 # Predict on the test set
22 y_pred = lin_reg.predict(X_test)
23
24 # Evaluate the model
25 mse = mean_squared_error(y_test, y_pred)
26 r2 = r2_score(y_test, y_pred)
27
28 print(f"Mean Squared Error: {mse}")
29 print(f"R^2 Score: {r2}")
30
31 # Plot the results

```

```
32 plt.scatter(X_test, y_test, color='blue', label='Actual')
33 plt.plot(X_test, y_pred, color='red', label='Predicted')
34 plt.xlabel("X")
35 plt.ylabel("y")
36 plt.title("Linear Regression")
37 plt.legend()
38 plt.show()
39
```

```
➦ Mean Squared Error: 1.0434333815695171
R-squared: 0.74344152333071327
```