

**Name:** Vaibhav Bansal **Registration Num.:** 229310063 **Class and Section:** 5E **Lab Experiment(s):** *from 1 till 7*

Experiment 1 : Data Cleaning & EDA

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings

data = pd.read_csv("/content/data.csv")

data.head(5)

{"summary":{"\n  \"name\": \"data\",\n  \"rows\": 10,\n  \"fields\": [\n    {\n      \"column\": \"Country\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"France\",\n          \"Spain\",\n          \"Germany\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"Age\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 7.693792591722527,\n          \"min\": 27.0,\n          \"max\": 50.0,\n          \"num_unique_values\": 9,\n          \"samples\": [\n            50.0,\n            27.0,\n            35.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n      {\n        \"column\": \"Salary\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 12265.579661982732,\n          \"min\": 48000.0,\n          \"max\": 83000.0,\n          \"num_unique_values\": 9,\n          \"samples\": [\n            83000.0,\n            48000.0,\n            52000.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n      {\n        \"column\": \"Purchased\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 2,\n          \"samples\": [\n            \"Yes\",\n            \"No\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      }\n    ]\n  }],\n  \"type\": \"dataframe\",\n  \"variable_name\": \"data\"}
```

data.shape

(10, 4)

```
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)

data.head(2)

{"summary":{"\n  \"name\": \"data\",\n  \"rows\": 10,\n  \"fields\": [\n    {\n      \"column\": \"Country\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"France\",\n          \"Spain\",\n          \"Germany\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"Age\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 7.693792591722527,\n          \"min\": 27.0,\n          \"max\": 50.0,\n          \"num_unique_values\": 9,\n          \"samples\": [\n            50.0,\n            27.0,\n            35.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n      {\n        \"column\": \"Salary\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 12265.579661982732,\n          \"min\": 48000.0,\n          \"max\": 83000.0,\n          \"num_unique_values\": 9,\n          \"samples\": [\n            83000.0,\n            48000.0,\n            52000.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n      {\n        \"column\": \"Purchased\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 2,\n          \"samples\": [\n            \"Yes\",\n            \"No\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      }\n    ]\n  }],\n  \"type\": \"dataframe\",\n  \"variable_name\": \"data\"}
```

data.tail(2)

```
{"summary":{"\n  \"name\": \"data\",\n  \"rows\": 2,\n  \"fields\": [\n    {\n      \"column\": \"Country\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"France\",\n          \"Germany\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"Age\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 9.192388155425117,\n          \"min\": 37.0,\n          \"max\": 50.0,\n          \"num_unique_values\": 2,\n          \"samples\": [\n            50.0,\n            37.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n      {\n        \"column\": \"Salary\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 11313.70849898476,\n          \"min\": 67000.0,\n          \"max\": 83000.0,\n          \"num_unique_values\": 2,\n          \"samples\": [\n            67000.0,\n            83000.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n      {\n        \"column\": \"Purchased\",\n        \"properties\": {\n          \"dtype\": \"string\",\n          \"num_unique_values\": 2,\n          \"samples\": [\n            \"Yes\",\n            \"No\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      }\n    ]\n  }],\n  \"type\": \"dataframe\",\n  \"variable_name\": \"data\"}
```

```
\ "num_unique_values\ ": 2,\n      \ "samples\ ": [\n          \ "Yes\ ",\n          \ "No\ "\n      ],\n      \ "semantic_type\ ": \ "\",\n      \ "description\ ": \ "\n      }\n      }\n      ]\n      }", "type": "dataframe"}
```


```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Country     10 non-null     object
1   Age         9 non-null      float64
2   Salary      9 non-null      float64
3   Purchased   10 non-null     object
dtypes: float64(2), object(2)
memory usage: 448.0+ bytes
```

```
data.isnull().sum()
```

```
Country      0
Age           1
Salary        1
Purchased     0
dtype: int64
```

```
plt.figure(figsize=(25,25))
sns.heatmap(data.isnull())
plt.show()
```



```
missing_value_percent = data.isnull().sum() / data.shape[0] * 100
print(missing_value_percent)

Country      0.0
Age          10.0
Salary       10.0
Purchased    0.0
dtype: float64

missing_value_column = missing_value_percent[missing_value_percent > 17].keys()
print(missing_value_column)

Index([], dtype='object')

data1 = data.drop(columns = missing_value_column)

data1.shape

(10, 4)

data2 = data1.dropna()

data2.shape

(8, 4)

plt.figure(figsize=(25,25))
sns.heatmap(data1.isnull())
plt.show()
```



```
data2.isnull().sum().sum()
```

```
0
```

Experiment 2: Outliers estimation & Plot it through violen/whisker

```
# Importing
import sklearn
from sklearn.datasets import load_diabetes
import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
diabetics = load_diabetes()

# Create the dataframe
column_name = diabetics.feature_names
df_diabetics = pd.DataFrame(diabetics.data)
df_diabetics.columns = column_name
print(df_diabetics.head())
```

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	

	s4	s5	s6
0	-0.002592	0.019907	-0.017646
1	-0.039493	-0.068332	-0.092204
2	-0.002592	0.002861	-0.025930
3	0.034309	0.022688	-0.009362
4	-0.002592	-0.031988	-0.046641

```
import seaborn as sns
import matplotlib.pyplot as plt



def removal_box_plot(df, column, threshold):
    sns.boxplot(df[column])
    plt.title(f'Original Box Plot of {column}')
    plt.show()

    removed_outliers = df[df[column] <= threshold]


    sns.boxplot(removed_outliers[column])
    plt.title(f'Box Plot without Outliers of {column}')
    plt.show()
    return removed_outliers

threshold_value = 0.12
```

```
no_outliers = removal_box_plot(df_diabetics, 'bmi', threshold_value)
```



```
fig, ax = plt.subplots(figsize=(6, 4))
ax.scatter(df_diabetics['bmi'], df_diabetics['bp'])
ax.set_xlabel('(body mass index of people)')
ax.set_ylabel('(bp of the people )')
plt.show()
```




```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

outlier_indices = np.where((df_diabetics['bmi'] > 0.12) &
                           (df_diabetics['bp'] < 0.8))

no_outliers = df_diabetics.drop(outlier_indices[0])

# Scatter plot without outliers
fig, ax_no_outliers = plt.subplots(figsize=(6, 4))
ax_no_outliers.scatter(no_outliers['bmi'], no_outliers['bp'])
ax_no_outliers.set_xlabel('(body mass index of people)')
ax_no_outliers.set_ylabel('(bp of the people )')
plt.show()
```





```
import numpy as np

threshold_z = 2

outlier_indices = np.where(z > threshold_z)[0]
no_outliers = df_diabetics.drop(outlier_indices)
print("Original DataFrame Shape:", df_diabetics.shape)
print("DataFrame Shape after Removing Outliers:", no_outliers.shape)

Original DataFrame Shape: (442, 10)
DataFrame Shape after Removing Outliers: (426, 10)

# IQR
Q1 = np.percentile(df_diabetics['bmi'], 25, method='midpoint')
Q3 = np.percentile(df_diabetics['bmi'], 75, method='midpoint')
IQR = Q3 - Q1
print(IQR)

0.06520763046978838

# Above Upper bound
upper = Q3+1.5*IQR
upper_array = np.array(df_diabetics['bmi'] >= upper)
print("Upper Bound:", upper)
print(upper_array.sum())

# Below Lower bound
lower = Q1-1.5*IQR
lower_array = np.array(df_diabetics['bmi'] <= lower)
print("Lower Bound:", lower)
print(lower_array.sum())
```

```

Upper Bound: 0.12879000811776306
3
Lower Bound: -0.13204051376139045
0

# Importing
import sklearn
from sklearn.datasets import load_diabetes
import pandas as pd

# Load the dataset
diabetes = load_diabetes()

# Create the dataframe
column_name = diabetes.feature_names
df_diabetes = pd.DataFrame(diabetes.data)
df_diabetes.columns = column_name
df_diabetes.head()
print("Old Shape: ", df_diabetes.shape)

''' Detection '''
# IQR
# Calculate the upper and lower limits
Q1 = df_diabetes['bmi'].quantile(0.25)
Q3 = df_diabetes['bmi'].quantile(0.75)
IQR = Q3 - Q1
lower = Q1 - 1.5*IQR
upper = Q3 + 1.5*IQR

# Create arrays of Boolean values indicating the outlier rows
upper_array = np.where(df_diabetes['bmi'] >= upper)[0]
lower_array = np.where(df_diabetes['bmi'] <= lower)[0]

# Removing the outliers
df_diabetes.drop(index=upper_array, inplace=True)
df_diabetes.drop(index=lower_array, inplace=True)

# Print the new shape of the DataFrame
print("New Shape: ", df_diabetes.shape)

Old Shape: (442, 10)
New Shape: (439, 10)

```

Experiment 3: Imputation of Missing Values KNN, Interpolation, Min-Max/Z-score.

```

# import necessary libraries
import numpy as np
import pandas as pd

# import the KNNImputer class
from sklearn.impute import KNNImputer

# create dataset for marks of a student

```

```
dict = {'Maths': [80, 90, np.nan, 95],
        'Chemistry': [60, 65, 56, np.nan],
        'Physics': [np.nan, 57, 80, 78],
        'Biology': [78, 83, 67, np.nan]}

# creating a data frame from the list
Before_imputation = pd.DataFrame(dict)
# print dataset before imputation
print("Data Before performing imputation\n", Before_imputation)

# create an object for KNNImputer
imputer = KNNImputer(n_neighbors=2)
After_imputation = imputer.fit_transform(Before_imputation)
# print dataset after performing the operation
print("\n\nAfter performing imputation\n", After_imputation)
```

```
Data Before performing imputation
   Maths  Chemistry  Physics  Biology
0   80.0         60.0     NaN    78.0
1   90.0         65.0     57.0    83.0
2    NaN         56.0     80.0    67.0
3   95.0         NaN    78.0     NaN
```

```
After performing imputation
[[80.  60.  68.5 78. ]
 [90.  65.  57.  83. ]
 [87.5 56.  80.  67. ]
 [95.  58.  78.  72.5]]
```

```
import pandas as pd
import numpy as np
a=pd.Series([0, 1, np.nan, 3,4,5,7])
a.interpolate()
```

```
0    0.0
1    1.0
2    2.0
3    3.0
4    4.0
5    5.0
6    7.0
dtype: float64
```

Experiment 4: Pivot table, Melt function.

```
import pandas as pd
import numpy as np
from sklearn.datasets import fetch_openml

X,y = fetch_openml("autos", version=1, as_frame=True, return_X_y=True)
data = X
data['target'] = y
```

```
pivot = np.round(pd.pivot_table(data, values='price',
                                index='num-of-doors',
                                columns='fuel-type',
                                aggfunc=np.mean),2)
```

pivot

<ipython-input-2-3c77819066fb>:1: FutureWarning: The default value of observed=False is deprecated and will change to observed=True in a future version of pandas. Specify observed=False to silence this warning and retain the current behavior

<ipython-input-2-3c77819066fb>:1: FutureWarning: The provided callable <function mean at 0x7f58e7b6c280> is currently using DataFrameGroupBy.mean. In a future version of pandas, the provided callable will be used directly. To keep current behavior pass the string "mean" instead.

```
pivot = np.round(pd.pivot_table(data, values='price',
```

```
{"summary":{"\n  \"name\": \"pivot\",\n  \"rows\": 2,\n  \"fields\": [\n    {\n      \"column\": \"num-of-doors\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"two\",\n          \"four\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"diesel\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1472.4650190072437,\n        \"min\": 14350.0,\n        \"max\": 16432.38,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          14350.0,\n          16432.38\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"column\": \"gas\",\n  \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 233.3805931306195,\n    \"min\": 12762.76,\n    \"max\": 13092.81,\n    \"num_unique_values\": 2,\n    \"samples\": [\n      12762.76,\n      13092.81\n    ],\n    \"semantic_type\": \"\",\n    \"description\": \"\"\n  }\n  ]\n},\"type\":\"dataframe\",\"variable_name\":\"pivot\"}
```

```
pivot = np.round(pd.pivot_table(data, values='price',
                                index=['num-of-doors', 'body-style'],
                                columns=['fuel-type', 'fuel-system'],
                                aggfunc=np.mean,
                                fill_value=0),2)
```

pivot

<ipython-input-3-14e7dc7c7159>:1: FutureWarning: The default value of observed=False is deprecated and will change to observed=True in a future version of pandas. Specify observed=False to silence this warning and retain the current behavior

<ipython-input-3-14e7dc7c7159>:1: FutureWarning: The provided callable <function mean at 0x7f58e7b6c280> is currently using DataFrameGroupBy.mean. In a future version of pandas, the provided callable will be used directly. To keep current behavior pass the string "mean" instead.

```
pivot = np.round(pd.pivot_table(data, values='price',
```

```
{"summary":{"\n  \"name\": \"pivot\",\n  \"rows\": 7,\n  \"fields\": [\n    {\n      \"column\": [\n        \"diesel\",\n        \"idi\"\n      ],\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 10517.712677237792,\n        \"min\": 0.0,\n        \"max\": 28176.0,\n        \"num_unique_values\": 6,\n        \"samples\": [\n          7788.0,\n          16328.92,\n          7437.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": [\n        \"gas\",\n        \"1bbl\"\n      ],\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 4163.206746769564,\n        \"min\": 0.0,\n        \"max\": 8811.67,\n        \"num_unique_values\": 4,\n        \"samples\": [\n          8811.67,\n          7054.43,\n          0.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": [\n        \"gas\",\n        \"2bbl\"\n      ],\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2943.290339369966,\n        \"min\": 0.0,\n        \"max\": 8249.0,\n        \"num_unique_values\": 7,\n        \"samples\": [\n          7813.71,\n          7711.19,\n          6701.67\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": [\n        \"gas\",\n        \"4bbl\"\n      ],\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 4590.378524697065,\n        \"min\": 0.0,\n        \"max\": 12145.0,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          12145.0,\n          0.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": [\n        \"gas\",\n        \"mfi\"\n      ],\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 4899.9314280916215,\n        \"min\": 0.0,\n        \"max\": 12964.0,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          12964.0,\n          0.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": [\n        \"gas\",\n        \"mpfi\"\n      ],\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 4749.15139588016,\n        \"min\": 10618.0,\n        \"max\": 23540.5,\n        \"num_unique_values\": 7,\n        \"samples\": [\n          10618.0,\n          18425.68\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": [\n        \"gas\",\n        \"spdi\"\n      ],\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 5104.217823594522,\n        \"min\": 0.0,\n        \"max\": 11479.43,\n        \"num_unique_values\": 7,\n        \"samples\": [\n          10618.0,\n          18425.68\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"column\": [\n    \"gas\",\n    \"mpfi\"\n  ],\n  \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 4749.15139588016,\n    \"min\": 10618.0,\n    \"max\": 23540.5,\n    \"num_unique_values\": 7,\n    \"samples\": [\n      10618.0,\n      18425.68\n    ],\n    \"semantic_type\": \"\",\n    \"description\": \"\"\n  }\n  ]\n},\"type\":\"dataframe\",\"variable_name\":\"pivot\"}
```



	ID	variable	value
0	1	Name	Tom
1	2	Name	Jerry
2	3	Name	Spike
3	1	Role	Cat
4	2	Role	Mouse
5	3	Role	Dog

```
#multiple columns as id_vars
df_melted = pd.melt(df, id_vars=["ID", "Name"], value_vars=["Role"])
print(df_melted)
```

	ID	Name	variable	value
0	1	Tom	Role	Cat
1	2	Jerry	Role	Mouse
2	3	Spike	Role	Dog

```
#skipping columns in melt function
df_melted = pd.melt(df, id_vars=["Name"], value_vars=["Role"])
print(df_melted)
```

	Name	variable	value
0	Tom	Role	Cat
1	Jerry	Role	Mouse
2	Spike	Role	Dog

Experiment 5: Image normalization (rescale -1,1).

```
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt

# Load an image
image_path = r"/content/mascot-logo-design_P1_900x420.jpg"

try:
    image = Image.open(image_path)
    image = image.convert('RGB') # Ensure image is in RGB format
except FileNotFoundError:
    print(f"File not found: {image_path}")
    raise
except Exception as e:
    print(f"Error occurred: {e}")
    raise

# Convert image to numpy array
image_array = np.array(image).astype('float32')

# Normalize the image to the range [-1, 1]
normalized_image_array = (image_array / 127.5) - 1

# Verify normalization
print(f"Min value in normalized image: {np.min(normalized_image_array)}")
print(f"Max value in normalized image: {np.max(normalized_image_array)}")
```



```

# Convert normalized image back to the range [0, 255] for visualization
denormalized_image_array = (normalized_image_array + 1) * 127.5
denormalized_image_array = denormalized_image_array.astype('uint8')

# Display the original and normalized images
fig, axes = plt.subplots(1, 2, figsize=(12, 6))

axes[0].imshow(image)
axes[0].set_title("Original Image")
axes[0].axis('off')

axes[1].imshow(denormalized_image_array)
axes[1].set_title("Normalized Image (rescaled to [0, 255] for display)")
axes[1].axis('off')

plt.show()

Min value in normalized image: -1.0
Max value in normalized image: 1.0

```

Original Image



Normalized Image (rescaled to [0, 255] for display)



Experiment 6: Linear regression and Logistic regression.

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Create a synthetic dataset for linear regression
np.random.seed(0)
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=0)

# Create and train the linear regression model

```

```

lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)

# Predict on the test set
y_pred = lin_reg.predict(X_test)


# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R^2 Score: {r2}")

# Plot the results
plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.plot(X_test, y_pred, color='red', label='Predicted')
plt.xlabel("X")
plt.ylabel("y")
plt.title("Linear Regression")
plt.legend()
plt.show()

```

Mean Squared Error: 1.0434333815695171  
R^2 Score: 0.7424452332071367



Experiment 7: Hypothesis testing (1 way ANOVA).



```
import numpy as np
from scipy.stats import f_oneway

# Generate sample data for three groups
np.random.seed(0)
group1 = np.random.normal(20, 5, 30)
group2 = np.random.normal(22, 5, 30)
group3 = np.random.normal(19, 5, 30)

# Perform one-way ANOVA
F_statistic, p_value = f_oneway(group1, group2, group3)

print(f"F-statistic: {F_statistic}")
print(f"P-value: {p_value}")

# Significance level
alpha = 0.05

# Decision
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference.")
else:
    print("Fail to reject the null hypothesis: There is no significant diff.")

F-statistic: 4.588157626753169
P-value: 0.012753331593793025
Reject the null hypothesis: There is a significant difference between the group means.
```